



**NANYANG  
TECHNOLOGICAL  
UNIVERSITY**

---

**SINGAPORE**

ISC17 SCC CODING CHALLENGE REPORT

SUBMITTED BY:  
NANYANG TECHNOLOGICAL UNIVERSITY

CODE CONTRIBUTOR(S):  
LIU SIYUAN

# Contents

<b>1</b>	<b>Optimizations</b>	<b>1</b>
1.1	Compilation optimization . . . . .	1
1.2	Code optimization . . . . .	1
1.2.1	Optimize algorithm in <code>add_vuspsi_k</code> . . . . .	1
1.2.2	Accelerate BLAS . . . . .	3
1.2.3	Accelerate (Sca)LAPACK . . . . .	3
1.2.4	Accelerate FFT . . . . .	4
1.2.5	Pinned memory for CUDA . . . . .	4
1.3	Runtime optimization . . . . .	4
1.3.1	Command line options . . . . .	4
1.3.2	Process binding . . . . .	4
<b>2</b>	<b>Compilation and run</b>	<b>5</b>
2.1	Fortran thunking interface . . . . .	10
2.2	Nvidia MPS . . . . .	10
<b>3</b>	<b>Result</b>	<b>10</b>

# 1 Optimizations

This section describes various optimization techniques team NTU employed to accelerate miniDFT.

## 1.1 Compilation optimization

We first optimized miniDFT by linking optimized libraries and using optimized compilers for our platform. In particular, we linked miniDFT against Intel Math Kernel Library (MKL) and used Intel compiler and Intel MPI to compile miniDFT. Also, with extensive testing, we discovered that MPI provides better parallelism than OpenMP and it is also more stable (OpenMP version gives wrong result sometimes). Therefore, we disabled OpenMP when compiling miniDFT (linking to `libmkl_sequential.a` and compile without `__OPENMP` and `-qopenmp`).

We also optimized compilation by tuning pre-processor macros. We discovered that there is a macro named `__NONBLOCKING_FFT` which controls what MPI routines are used for the `fft_scatter` function. The non-blocking version with asynchronous send and receive appears to perform much better on InfiniBand interconnect than the all to all version.

Below are the options we used in the Makefile for compilation optimization.

```
1 DFLAGS = -D__INTEL -D__FFTW -D_SCALAPACK -D__MPI -D__NONBLOCKING_FFT
2
3 MATH_LIBS = ${MKLR00T}/lib/intel64/libmkl_scalapack_lp64.a \
4             -Wl,--start-group \
5             ${MKLR00T}/lib/intel64/libmkl_intel_lp64.a \
6             ${MKLR00T}/lib/intel64/libmkl_sequential.a \
7             ${MKLR00T}/lib/intel64/libmkl_core.a \
8             ${MKLR00T}/lib/intel64/libmkl_blacs_intelmpi_lp64.a \
9             -Wl,--end-group -lpthread -lm -ldl
10
11 CC = mpiicc
12 CFLAGS = -O3 -I./ -I${MKLR00T}/include
13
14 FC = mpiifort
15 FFLAGS = -O3 -fpp -I./ -I${MKLR00T}/include
16
17 LD = mpiifort
18 LDFLAGS =
```

## 1.2 Code optimization

During our preparation, we tried to optimize the code by changing the algorithms and porting part of the code to CUDA. Below are the detailed descriptions to our code optimization.

### 1.2.1 Optimize algorithm in `add_vuspsi_k`

In the original code, `add_vuspsi_k` calls `calbec` to perform ZGEMV matrix vector multiplication calculations and do a sum reduce over the resulting vector (as shown below).

```

1 DO ibnd = 1, m
2     ! JRD: Compute becp for just this ibnd here
3     !
4     CALL calbec ( n, vkb, psi, becp, ibnd )
5     !write(*,*) 'Computing becp', ibnd
6     !
7     ! other code ...
8     !
9 END DO

1 !-----
2 SUBROUTINE calbec_k ( npw, beta, psi, betapsi, ibnd )
3 !-----
4     !
5     ! other code ...
6     !
7     CALL ZGEMV( 'C', npw, nkb, (1.0_DP,0.0_DP), beta, npwx, psi(:,ibnd), 1, &
8         (0.0_DP, 0.0_DP), betapsi, 1 )
9     CALL mp_sum( betapsi( : ), intra_bgrp_comm )
10    !
11    CALL stop_clock( 'calbec' )
12    !
13    RETURN
14    !
15 END SUBROUTINE calbec_k

```

We noticed that the loop in `add_vuspsi_k` is actually calling `calbec` to perform ZGEMV operations over the matrix `psi`. Each iteration in the loop corresponds to one column in the matrix. This is equivalent to performing a single ZGEMM matrix matrix multiplication on `vkb` and `psi`. As ZGEMV is memory-bound and ZGEMM is compute-bound, we obviously prefer to use ZGEMM.

In our optimized version of `add_vuspsi_k`, it performs a single ZGEMM operation instead of calling `calbec` in the loop. And then it performs the sum reduce over the resulting matrix (instead of vector). This reduced 8488 calls to ZGEMV to 5 calls to ZGEMM for the input `pe-23.LOCAL.in`.

```

1 USE mp_global, ONLY : intra_bgrp_comm
2 USE mp,          ONLY : mp_sum
3 USE wvfct,       ONLY : npwx
4 !
5 ! other code ...
6 !
7 COMPLEX(DP), ALLOCATABLE :: betapsi(:, :)
8 !
9 ! other code ...
10 !
11 ALLOCATE( betapsi(SIZE(vkb,2), m) )
12 betapsi(:, :) = (0.0_DP, 0.0_DP)
13 CALL ZGEMM('C', 'N', SIZE(vkb,2), m, n, (1.0_DP, 0.0_DP), vkb, npwx, &
14     psi, lda, (0.0_DP, 0.0_DP), betapsi, SIZE(vkb,2))
15 CALL mp_sum(betapsi(:, :), intra_bgrp_comm)
16
17 DO ibnd = 1, m
18     ! nested loop ...

```

```

19     ps(ikb,ibnd) = ps(ikb,ibnd) + &
20         deeq(ih,jh,na,current_spin) * betapsi(jkb,ibnd)
21 END DO

```

### 1.2.2 Accelerate BLAS

As mentioned in Section 1.2.1, ?GEMMs are compute-bound routines. Therefore, they are perfect targets for offloading computation to GPUs. In our case, we used the Fortran thunking interface to cuBLAS provided by Nvidia to offload zGEMM to GPUs. We decided to offload calls to zGEMM in `cegterg.f90` and `add_vuspsi.f90` (after our optimization in Section 1.2.1). Some macros help us easily switch between MKL BLAS and cuBLAS.

```

1 #if defined(__CUDA) && defined(__CUBLAS)
2 #define ZGEMM cublas_ZGEMM
3 #endif

```

As we disabled OpenMP, we run one MPI rank per CPU core. This means that we would launch a huge number of processes on a single GPU. As they all have different CUDA context, the calls to GPU get serialized between those processes. In order to overcome this problem, we used Nvidia Multi-Process Service (MPS). Nvidia MPS allows multiple processes to share a CUDA context on the same GPU. The efficiency of zGEMM are greatly improved on the GPU after using Nvidia MPS.

### 1.2.3 Accelerate (Sca)LAPACK

In the original code, ScaLAPACK is used in `cdiaghg.f90` to perform the diagonalization. We decided to port the diagonalization to GPU. As we are not aware of any mature ScaLAPACK implementation on the GPU, we used MAGMA, a serial hybrid (CPU + GPU) LAPACK library.

Because we used a serial LAPACK library, we have to create a serial version of `pcdiaghg` (named `cdiaghg_gpu`). We borrowed some code from the `qe-gpu-plugin` GitHub repository and fixed an error caused by the magic number inside. After the bug fix, the diagonalization could be successfully performed on the GPU with MAGMA. For each band group, only the root rank performs the diagonalization (hence making it serial). We used the single GPU interface (`magmaf_zhegvd`) instead of the multi-GPU interface because we realized that the single GPU version performs better for the matrix size we have to deal with.

```

1 IF ( me_bgrp == root_bgrp ) THEN
2     CALL magmaf_init()
3     !
4     ! other code ...
5     CALL magmaf_zhegvd(1, 'V', 'U', n, v, ldh, s, ldh, e, &
6                         work, lwork, rwork, lrwork, iwork, liwork, info)
7     !
8     ! other code ...
9     !
10    CALL magmaf_zhegvx( 1, 'V', 'I', 'U', n, h, ldh, s, ldh, &
11                        0.D0, 0.D0, 1, m, abstol, mm, e, v, ldh, &
12                        work, lwork, rwork, iwork, ifail, info )
13    !

```

```

14      ! other code ...
15      !
16      CALL magma_finalize()
17 END IF
18 !
19 ! ... broadcast eigenvectors and eigenvalues to all other processors
20 !
21 CALL mp_bcast( e, root_bgrp, intra_bgrp_comm )
22 CALL mp_bcast( v, root_bgrp, intra_bgrp_comm )

```

### 1.2.4 Accelerate FFT

We tried to also port FFT to GPU with the cuFFT library. To make the porting process easy, we decided to use the PGI compiler (CUDA Fortran). We successfully ported all FFT operations to GPU, including the task group version, with CUDA-aware MPI. We also enabled GPUDirect RDMA on our cluster nodes to accelerate the communication as `fft_scatter` takes quite some time. In the end, however, we disabled our optimization on FFT because it appears to be slower than the CPU (88 CPU cores v.s. 8 Tesla P100s). Maybe it's because we have so many CPU cores and cuFFT cannot provide too much speed up compared to MKL FFT.

Interested readers could refer to `vloc_psi_gpu.f90` and other FFT related files with the suffix `_gpu.f90`.

### 1.2.5 Pinned memory for CUDA

Pinned memory is an optimization technique for accelerating data transfers between CPU and GPU. It allocates page-locked host memory to speed up the transfer. In our code, if PGI compiler is used, the large arrays in `pcegtarg` are allocated as pinned memory. This optimization is not available if Intel compiler is used. In practice, we cannot notice a difference as the array are quite large and the benefit of pinned memory becomes negligible.

## 1.3 Runtime optimization

### 1.3.1 Command line options

We briefly studied the background of miniDFT to understand how the various command line options would impact the performance. With what we learned and our testing, it seemed that none of the command line options need to be adjusted. Optimizations like band group parallelism and task group parallelism seem to be beneficial only when running miniDFT at a very large scale (e.g. thousands of cores). If we enable those options manually, the performance will actually drop on our platform (small scale).

### 1.3.2 Process binding

As we ported part of the code to GPU, miniDFT now runs on a heterogeneous platform. Therefore, it's necessary to bind MPI processes to the CPU socket which is directly connected to the particular GPU the process controls in order to improve performance. We used `numactl` and bind processes according to their rank when launching the job (we have 4 GPUs in a single node).

```

1 #!/bin/bash
2
3 rank=$(( ${PMI_RANK} % 4 ))
4 case ${rank} in
5 [0]) numactl --cpunodebind=0 ./mini_dft -in pe-23.LOCAL.in;;
6 [1]) numactl --cpunodebind=0 ./mini_dft -in pe-23.LOCAL.in;;
7 [2]) numactl --cpunodebind=1 ./mini_dft -in pe-23.LOCAL.in;;
8 [3]) numactl --cpunodebind=1 ./mini_dft -in pe-23.LOCAL.in;;
9 esac

```

## 2 Compilation and run

This section is meant to guide the reader to successfully compile and run our optimized version of miniDFT. First, please make sure the following pre-requisites are met. For compiling MAGMA, please refer to their documentation<sup>1</sup>. It's well documented and not hard to compile. For Fortran thunking interface and Nvidia MPS<sup>2</sup>, please refer to Section 2.1 and 2.2.

- MAGMA (compiled with Intel compiler and MKL, without OpenMP)
- Intel compilers and Intel MPI
- GCC
- Intel MKL
- CUDA (with Fortran thunking cuBLAS interface)
- Nvidia MPS server (running on the node)
- numactl

To compile the code, use the following Makefile (available as a gist<sup>3</sup>) and type `make`. Make sure you have `MKLROOT` and `MAGMAROOT` as environment variables. Also, make sure you have `fortran_thunking.o`.

```

1 # macros
2 DFLAGS = -D__INTEL -D__FFTW -D__SCALAPACK -D__MPI -D__CUDA -D__MAGMA -D__ZHEGVD -
   D__CUBLAS -D__NONBLOCKING_FFT
3
4 # libraries
5 MATH_LIBS = ${MKLROOT}/lib/intel64/libmkl_scalapack_lp64.a \
6             -Wl,--start-group \
7             ${MKLROOT}/lib/intel64/libmkl_intel_lp64.a \
8             ${MKLROOT}/lib/intel64/libmkl_sequential.a \
9             ${MKLROOT}/lib/intel64/libmkl_core.a \
10            ${MKLROOT}/lib/intel64/libmkl_blacs_intelmpi_lp64.a \
11            -Wl,--end-group -lpthread -lm -ldl
12

```

<sup>1</sup><http://icl.cs.utk.edu/projectsfiles/magma/dxygen/installing.html>

<sup>2</sup>[https://docs.nvidia.com/deploy/pdf/CUDA\\_Multi.Process.Service.Overview.pdf](https://docs.nvidia.com/deploy/pdf/CUDA_Multi.Process.Service.Overview.pdf)

<sup>3</sup><https://gist.github.com/koallen/da937a36b2f292e86a6e872e78354789>

```

13 LIBS = ${MAGMAROOT}/lib/libmagma.a -lcuda -lcudart -lcublas -lcusparse $(
    MATH_LIBS) -lstdc++
14
15 CC = mpiicc
16 CFLAGS = -O3 -I./ -I${MKLR00T}/include
17
18 FC = mpiifort
19 FFLAGS = -O3 -fpp -I./ -I${MKLR00T}/include
20
21 LD = mpiifort
22 LDFLAGS =
23
24 CFLAGS += $(DFFLAGS)
25
26 FFLAGS += $(DFFLAGS)
27
28 OBJECTS += \
29     gpu_aux.o \
30     kind.o \
31     constants.o \
32     radial_grids.o \
33     atom.o \
34     io_global.o \
35     parallel_include.o \
36     mp.o \
37     mp_global.o \
38     cuda_global.o \
39     parser.o \
40     parameters.o \
41     input_parameters.o \
42     io_files.o \
43     control_flags.o \
44     cell_base.o \
45     check_stop.o \
46     clocks.o \
47     fft_types.o \
48     fft_base.o \
49     fft_base_gpu.o \
50     random_numbers.o \
51     ions_base.o \
52     descriptors.o \
53     electrons_base.o \
54     version.o \
55     environment.o \
56     error_handler.o \
57     cufft.o \
58     fft_scalar.o \
59     fft_scalar_gpu.o \
60     fft_custom.o \
61     recvec.o \
62     pwcom.o \
63     stick_base.o \
64     stick_set.o \

```



```
65     data_structure_custom.o \
66     fft_parallel.o \
67     fft_parallel_gpu.o \
68     fft_interfaces.o \
69     fft_interfaces_gpu.o \
70     wrappers.o \
71     funct.o \
72     griddim.o \
73     image_io_routines.o \
74     allocate_fft_custom.o \
75     ruotaijk.o \
76     xk_wk_collect.o \
77     mp_base.o \
78     mp_wave.o \
79     mp_image_global_module.o \
80     pseudo_types.o \
81     zhpev_drv.o \
82     ptoolkit.o \
83     read_cards.o \
84     read_namelists.o \
85     uspp.o \
86     upf_nml.o \
87     upf.o \
88     upf_to_internal.o \
89     read_pseudo.o \
90     recvec_subs.o \
91     run_info.o \
92     set_signal.o \
93     splinelib.o \
94     wavefunctions.o \
95     \
96     c_mkdir.o \
97     cptimer.o \
98     customize_signals.o \
99     eval_infix.o \
100    fft_stick.o \
101    md5.o \
102    md5_from_file.o \
103    memstat.o \
104    stack.o \
105    \
106    atomic_number.o \
107    capital.o \
108    cryst_to_car.o \
109    date_and_tim.o \
110    distools.o \
111    erf.o \
112    find_free_unit.o \
113    flush_unit.o \
114    functionals.o \
115    inpfiler.o \
116    int_to_char.o \
117    invmat.o \
```

```
118     latgen.o \
119     lsda_functionals.o \
120     matches.o \
121     recipso.o \
122     remove_tot_torque.o \
123     rgen.o \
124     simpsn.o \
125     sort.o \
126     sph_bes.o \
127     trimcheck.o \
128     volume.o \
129     ylmr2.o \
130     \
131     symm_base.o \
132     start_k.o \
133     scf_mod.o \
134     a2fmod.o \
135     buffers.o \
136     becmod.o \
137     add_vuspsi.o \
138     allocate_fft.o \
139     allocate_locpot.o \
140     allocate_nlpot.o \
141     allocate_wfc.o \
142     atomic_rho.o \
143     atomic_wfc.o \
144     g_psi_mod.o \
145     c_bands.o \
146     ccgdiag.o \
147     cdiag.o \
148     cdiag_gpu.o \
149     cegterg.o \
150     symme.o \
151     close_files.o \
152     coset.o \
153     data_structure.o \
154     deriv_drhoc.o \
155     divide.o \
156     divide_et_impera.o \
157     drhoc.o \
158     dvloc_of_g.o \
159     compute_deff.o \
160     newd.o \
161     coulomb_vcut.o \
162     exx.o \
163     clean_pw.o \
164     input.o \
165     electrons.o \
166     eqvect.o \
167     ewald.o \
168     g2_kin.o \
169     g_psi.o \
170     gk_sort.o \
```

```

171     gradcorr.o \
172     h_1psi.o \
173     h_psi.o \
174     hinit0.o \
175     init_at_1.o \
176     openfil.o \
177     init_run.o \
178     init_us_1.o \
179     init_us_2.o \
180     init_vloc.o \
181     interpolate.o \
182     irrek.o \
183     iweights.o \
184     kpoint_grid.o \
185     lchk_tauk.o \
186     memory_report.o \
187     mix_rho.o \
188     multable.o \
189     n_plane_waves.o \
190     para.o \
191     potinit.o \
192     print_clock_pw.o \
193     print_ks_energies.o \
194     read_input.o \
195     pwscf.o \
196     remove_atomic_rho.o \
197     rotate_wfc.o \
198     rotate_wfc_k.o \
199     s_1psi.o \
200     s_psi.o \
201     set_kup_and_kdw.o \
202     set_rhoc.o \
203     set_vrs.o \
204     setlocal.o \
205     setup.o \
206     stop_run.o \
207     struct_fact.o \
208     sum_band.o \
209     summary.o \
210     usnldiag.o \
211     v_of_rho.o \
212     vloc_of_g.o \
213     vloc_psi.o \
214     vloc_psi_gpu.o \
215     weights.o \
216     wfcinit.o
217
218 # rules
219 all: mini_dft $(if $(USE_HPCTK), mini_dft.hpcstruct)
220
221 %.o : %.f90
222     $(FC) $(FFLAGS) -c $<
223

```

```

224 mini_dft: $(OBJECTS)
225     $(LD) $(LDFLAGS) -o $@ fortran_thunking.o $(OBJECTS) $(LIBS) $(IPM)
226     -rm -f mini_dft.hpcstruct
227
228 mini_dft.hpcstruct: mini_dft
229     hpcstruct $<
230
231 .PHONY = clean
232 clean:
233     -rm -f *~ *.o *.mod mini_dft
234
235 .PHONY = tarball
236 tarball:
237     tar -czf mini_dft.tar.gz Makefile *.UPF *.UPF.nml *.h *.c *.f *.f90

```

After successful compilation, launch the program with MPI and the run script in Section 1.3.2 (adjust according to your environment). The number of MPI ranks should equal the number of CPU cores. For example,

```

1 $ mpirun -np 88 -ppn 44 -hosts compute0,compute1 bash run.sh

```

## 2.1 Fortran thunking interface

Nvidia provides the Fortran thunking cuBLAS interface. It's available in `/path/to/cuda/install/dir/src`. To compile it, do

```

1 $ nvcc -O3 -c fortran_thunking.c -Xcompiler -DCUBLAS_GFORTRAN

```

## 2.2 Nvidia MPS

To run Nvidia MPS, first set GPU to `EXCLUSIVE_PROCESS` mode. Then start the MPS server as root.

```

1 $ sudo nvidia-smi -c 3
2 $ sudo nvidia-cuda-mps-control -d

```

## 3 Result

We performed extensive benchmarking on our cluster for miniDFT. The specification of our nodes is in Table 1, and we have two identical nodes in total. The runtime we recorded for original and final versions are in Table 2. No command line option besides `-in pe-23.LOCAL.in` is used. The run script in Section 1.3.2 is used for launching jobs for the final version.

As shown in Table 2, we were able to achieve a speed up of 3.1 on our cluster for the input `pe-23.LOCAL.in`.

Table 1: Machine specification

<b>Item</b>	<b>Count</b>
E5-2699 v4 (2.2GHz, 22 cores) CPU	2
DDR4 16GB 2400MHz RAM	16 (256GB total)
Mellanox InfiniBand EDR adapter	1
Nvidia Tesla P100 accelerator	4
SAMSUNG 500GB SSD	1

Table 2: Runtime

<b>Version</b>	<b>Benchmark Time</b>
Original (only compilation optimization)	504.36 s
Final (comp ops, algo ops, cuBLAS, MAGMA, proc binding)	162.60 s